Clustering

Boston University CS 506 - Lance Galletti

What is a Clustering

A clustering is a grouping / assignment of objects (data points) such that objects in the same group / cluster are:
- similar to one another
- dissimilar to objects in other groups

# The Clustering Problem

Given a collection of data points

Find a clustering such that:
- **Similar** data points are in the **same cluster**
- **Dissimilar** data points are in **different clusters**

Questions:
- What does **similar** mean?
- How do we find a **clustering**?
- How do we know if we have found a **good clustering**?

Clusters can be Ambiguous

Types of Clusterings

Partitional
Each object belongs to exactly one cluster

Hierarchical
A set of nested clusters organized in a tree

Density-Based
Defined based on the local density of points

Soft Clustering
Each point is assigned to every cluster with a certain probability

Partitional Clustering

Partitional Clustering

Given **n** data points and a number **k** of clusters: partition the **n** data points into **k** clusters.

Suppose we are given all possible ways of distributing these **n** data points into these **k** buckets / clusters. How would we find the best such partition?

Recall our goal: **similar** items should belong to the **same cluster** & **dissimilar** items should belong to **different clusters**.

A good partition is one where the total dissimilarity of points within each cluster is small.

Example

VS

Clearly the clustering on the left has smaller intra-cluster distances than the one on the right. That is:

$$\sum_{k}^{K} \sum_{x_i, x_j \in C_k} d(x_i, x_j)$$

Is a smaller quantity

Example

VS

Given a distance function **d**, we can find points (not necessarily part of our dataset) for each cluster called **centroids** that are at the center of each cluster.

Example

VS

Q: When **d** is Euclidean, what is the **centroid** (also called **center of mass**) of **m** points $\{x_1, \dots, x_m\}$?

A: The mean / average of the points

Example

VS

Turns out when **d** is Euclidean:

$$\sum_{k}^{K} \sum_{x_i, x_j \in C_k} d(x_i, x_j)^2 = \sum_{k}^{K} |C_k| \sum_{x_i \in C_k} d(x_i, \mu_k)^2$$

K-means

Given **X = {x$_1$, ... , x$_n$}** our dataset and **k**

Find **k** points **{μ$_1$, ... , μ$_k$}** that minimize the **cost function**:

$$\sum_{i}^{k} \sum_{x \in C_i} \|x - \mu_i\|_2^2$$

When **k=1** and **k=n** this is easy. Why?

When **x$_i$** lives in more than 2 dimensions, this is a very difficult (**NP-hard**) problem

K-means - Lloyd's Algorithm

1. Randomly pick **k** centers $\{\mu_1, \ldots, \mu_k\}$
2. Assign each point in the dataset to its closest center
3. Compute the new centers as the means of each cluster
4. Repeat 2 & 3 until convergence

K-means - Lloyd's Algorithm
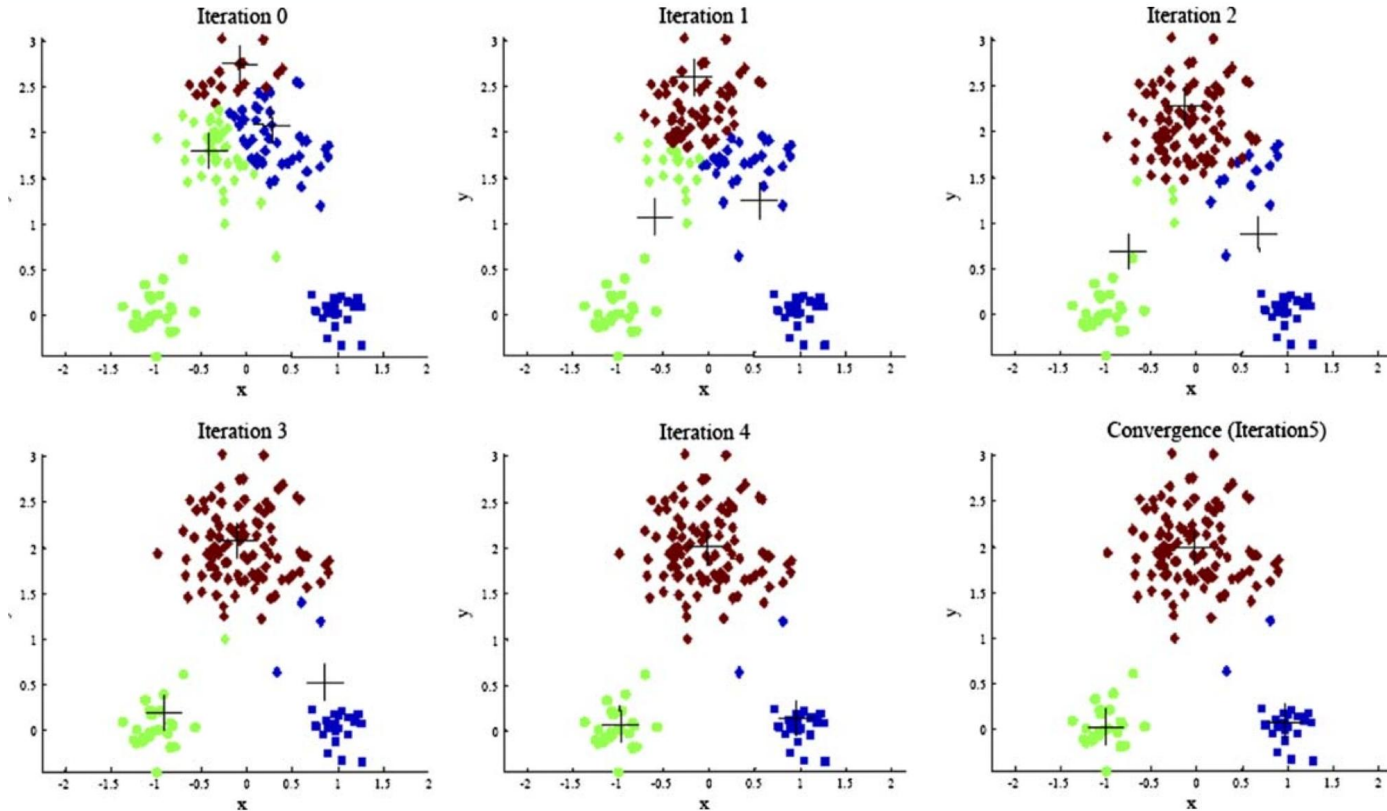
K-means - Lloyd's Algorithm

Will this algorithm always converge?

Proof (by contradiction): Suppose it does not converge. Then, either:

1. The minimum of the cost function is only reached in the limit (i.e. after an infinite number of iterations).
    Impossible because we are iterating over a finite set of partitions

2. The algorithm gets stuck in a cycle / loop
    Impossible since this would require having a clustering that has a lower cost than itself and we know:
    -   If old ≠ new clustering then the cost has improved
    -   If old = new clustering then the cost is unchanged

Conclusion: Lloyd's Algorithm always converges!

K-means - Lloyd's Algorithm

Will this always converge to the optimal solution?

K-means - Lloyd's Algorithm

The choice of initial points has a large influence on the resulting clustering

K-means - Initialization

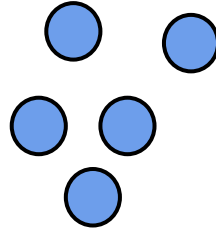One solution: Run Lloyd's algorithm multiple times and choose the result with the lowest cost.

This can still lead to bad results because of randomness.

Another solution: Try different initialization methods
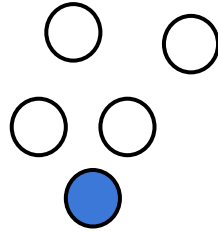
K-means - Random
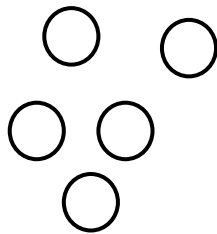
K-means - Random

Starting with initialization points too close to each other may problematic
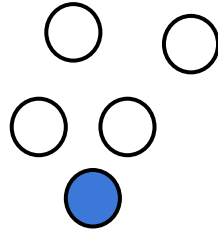
K-means - Farthest First Traversal

# K-means - Farthest First Traversal
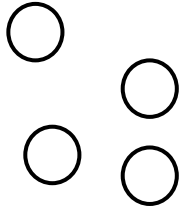
Pick the first center at random

K-means - Farthest First Traversal

Pick the next center to be the
point farthest from all previous

K-means - Farthest First Traversal

Pick the next center to be the
point farthest from all previous

# K-means - Farthest First Traversal

# K-means - FFT and outliers

# K-means - FFT and outliers

K-means - FFT and outliers

K-means - FFT and outliers

K-means - FFT and outliers

Random might have worked better here

K-means++

Initialize with a combination of the two methods:

1. Start with a random center
2. Let **D(x)** be the distance between **x** and the centers selected so far. Choose the next center with probability proportional to $D(x)^a$

When:
- **a = 0** : random initialization (all points have equal probability)
- **a = ∞** : farthest first traversal
- **a = 2** : K-means++

K-means++

K-means++

No reason to use k-means over k-means++

K-means++

Suppose we are given a black box that will generate a uniform random number between 0 and any **N**. How can we use this black box to select points with probability proportional to $D(x)^a$?

K-means++

Suppose we are given a black box that will generate a uniform random number between 0 and any **N**. How can we use this black box to select points with probability proportional to **D(x)^a**?

K-means++

Suppose we are given a black box that will generate a uniform random number between 0 and any **N**. How can we use this black box to select points with probability proportional to **D(x)²**?

Let's set **a = 2**

$$D(x)^2 = 3^2 = 9$$
$$D(y)^2 = 2^2 = 4$$
$$D(z)^2 = 1^2 = 1$$

K-means++

Suppose we are given a black box that will generate a uniform random number between 0 and any **N**. How can we use this black box to select points with probability proportional to **D(x)²**?

Let's set **a = 2**

$D(x)^2 = 3^2 = 9$
$D(y)^2 = 2^2 = 4$
$D(z)^2 = 1^2 = 1$

x

y   z

0   N

K-means++

Suppose we are given a black box that will generate a uniform random number between 0 and any **N**. How can we use this black box to select points with probability proportional to **D(x)^2**?

Let's set **a = 2**

$D(x)^2 = 3^2 = 9$
$D(y)^2 = 2^2 = 4$
$D(z)^2 = 1^2 = 1$

K-means++

Suppose we are given a black box that will generate a uniform random number between 0 and any **N**. How can we use this black box to select points with probability proportional to **D(x)²**?

Let's set **a = 2**

$D(x)^2 = 3^2 = 9$
$D(y)^2 = 2^2 = 4$
$D(z)^2 = 1^2 = 1$

K-means++

Suppose we are given a black box that will generate a uniform random number between 0 and any **N**. How can we use this black box to select points with probability proportional to **D(x)²**?

Let's set **a = 2**

$$D(x)^2 = 3^2 = 9$$
$$D(y)^2 = 2^2 = 4$$
$$D(z)^2 = 1^2 = 1$$

K-means++

Suppose we are given a black box that will generate a uniform random number between 0 and any **N**. How can we use this black box to select points with probability proportional to **D(x)²**?

Let's set **a = 2**

$D(x)^2 = 3^2 = 9$
$D(y)^2 = 2^2 = 4$
$D(z)^2 = 1^2 = 1$

N
= D(x)² + D(y)²
+ D(z)² = 14

K-means++

Suppose we are given a black box that will generate a uniform random number between 0 and any **N**. How can we use this black box to select points with probability proportional to **D(x)²**?

Using the black box, we can generate a number between 0 and N to determine which point to pick next. It will be chosen with probability proportional to **D(x)²**.

0

N
= $D(x)^2 + D(y)^2 + D(z)^2 = 14$

K-means++

What happens if the black box can only generate numbers between 0 and 1?

K-means - Limitations

K-means - Limitations

K-means - Limitations

How to choose the right k?

1. Iterate through different values of k (elbow method)
2. Use empirical / domain-specific knowledge
   Example: Is there a known approximate distribution of the data? (K-means is good for spherical gaussians)

K-means Variations

1. K-medians (uses the L_1 norm / manhattan distance)
2. K-medoids (any distance function + the centers must be in the dataset)
3. Weighted K-means (each point has a different weight when computing the mean)